

```
# This Python 3 environment comes with many helpful analytics
libraries installed
# It is defined by the kaggle/python Docker image:
https://github.com/kaggle/docker-python
# For example, here's several helpful packages to load

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the read-only "../input/"
directory
# For example, running this (by clicking run or pressing Shift+Enter)
will list all files under the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/)
that gets preserved as output when you create a version using "Save &
Run All"
# You can also write temporary files to /kaggle/temp/, but they won't
be saved outside of the current session
```

Customer Segmentation

In this project, I will be performing an unsupervised clustering of data on the customer's records from a groceries firm's database. Customer segmentation is the practice of separating customers into groups that reflect similarities among customers in each cluster. I will divide customers into segments to optimize the significance of each customer to the business. To modify products according to distinct needs and behaviours of the customers. It also helps the business to cater to the concerns of different types of customers.

TABLE OF CONTENTS

- [1. IMPORTING LIBRARIES](#)
- [2. LOADING DATA](#)
- [3. DATA CLEANING](#)
- [4. DATA PREPROCESSING](#)
- [5. DIMENSIONALITY REDUCTION](#)
- [6. CLUSTERING](#)
- [7. EVALUATING MODELS](#)

- 8. PROFILING
- 9. CONCLUSION
- 10. END

IMPORTING LIBRARIES

```
#Importing the Libraries
import numpy as np
import pandas as pd
import datetime
import matplotlib
import matplotlib.pyplot as plt
from matplotlib import colors
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from yellowbrick.cluster import KElbowVisualizer
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt, numpy as np
from mpl_toolkits.mplot3d import Axes3D
from sklearn.cluster import AgglomerativeClustering
from matplotlib.colors import ListedColormap
from sklearn import metrics
import warnings
import sys
if not sys.warnoptions:
    warnings.simplefilter("ignore")
np.random.seed(42)
```

LOADING DATA

```
# Loading the dataset with proper delimiter
file_path = "C:/Users/Manas/Downloads/marketing_campaign.csv"
df = pd.read_csv(file_path, sep='\t') # Assuming the data is tab-separated
```

For more information on the attributes visit [here](#).

DATA CLEANING

In this section

- Data Cleaning
- Feature Engineering

In order to, get a full grasp of what steps should I be taking to clean the dataset. Let us have a look at the information in data.

```
# Information on features
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2240 entries, 0 to 2239
Data columns (total 29 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   ID                                     2240 non-null   int64
1   Year_Birth                           2240 non-null   int64
2   Education                             2240 non-null   object
3   Marital_Status                       2240 non-null   object
4   Income                               2216 non-null   float64
5   Kidhome                              2240 non-null   int64
6   Teenhome                             2240 non-null   int64
7   Dt_Customer                          2240 non-null   object
8   Recency                              2240 non-null   int64
9   MntWines                             2240 non-null   int64
10  MntFruits                             2240 non-null   int64
11  MntMeatProducts                       2240 non-null   int64
12  MntFishProducts                       2240 non-null   int64
13  MntSweetProducts                      2240 non-null   int64
14  MntGoldProds                          2240 non-null   int64
15  NumDealsPurchases                     2240 non-null   int64
16  NumWebPurchases                       2240 non-null   int64
17  NumCatalogPurchases                   2240 non-null   int64
18  NumStorePurchases                     2240 non-null   int64
19  NumWebVisitsMonth                     2240 non-null   int64
20  AcceptedCmp3                          2240 non-null   int64
21  AcceptedCmp4                          2240 non-null   int64
22  AcceptedCmp5                          2240 non-null   int64
23  AcceptedCmp1                          2240 non-null   int64
24  AcceptedCmp2                          2240 non-null   int64
25  Complain                              2240 non-null   int64
26  Z_CostContact                         2240 non-null   int64
27  Z_Revenue                             2240 non-null   int64
28  Response                              2240 non-null   int64
dtypes: float64(1), int64(25), object(3)
memory usage: 507.6+ KB
```

From the above output, we can conclude and note that:

- There are missing values in income
- Dt_Customer that indicates the date a customer joined the database is not parsed as DateTime
- There are some categorical features in our data frame; as there are some features in dtype: object). So we will need to encode them into numeric forms later.

First of all, for the missing values, I am simply going to drop the rows that have missing income values.

```
# To remove the NA values
df = df.dropna()
print("The total number of data-points after removing the rows with
missing values are:", len(df))
```

The total number of data-points after removing the rows with missing values are: 2216

In the next step, I am going to create a feature out of "Dt_Customer" that indicates the number of days a customer is registered in the firm's database. However, in order to keep it simple, I am taking this value relative to the most recent customer in the record.

Thus to get the values I must check the newest and oldest recorded dates.

```
# Convert 'Dt_Customer' column to datetime
df["Dt_Customer"] = pd.to_datetime(df["Dt_Customer"])

# Extract dates
dates = []
for i in df["Dt_Customer"]:
    i = i.date()
    dates.append(i)

# Dates of the newest and oldest recorded customer
print("The newest customer's enrollment date in the records:",
max(dates))
print("The oldest customer's enrollment date in the records:",
min(dates))
```

The newest customer's enrollment date in the records: 2014-12-06
The oldest customer's enrollment date in the records: 2012-01-08

Creating a feature ("Customer_For") of the number of days the customers started to shop in the store relative to the last recorded date

```
# Created a feature "Customer_For"
days = []
d1 = max(dates) # Taking it to be the newest customer
for i in dates:
```

```

    delta = d1 - i
    days.append(delta.days) # Extracting only the number of days
df["Customer_For"] = days
df["Customer_For"] = pd.to_numeric(df["Customer_For"],
errors="coerce") # Converting to numeric

```

```

# Printing the first few rows to verify the changes
print(df.head())

```

	ID	Year_Birth	Education	Marital_Status	Income	Kidhome	Teenhome
0	5524	1957	Graduation	Single	58138.0	0	0
1	2174	1954	Graduation	Single	46344.0	1	1
2	4141	1965	Graduation	Together	71613.0	0	0
3	6182	1984	Graduation	Together	26646.0	1	0
4	5324	1981	PhD	Married	58293.0	1	0

	Dt_Customer	Recency	MntWines	...	AcceptedCmp3	AcceptedCmp4	\
0	2012-04-09	58	635	...	0	0	
1	2014-08-03	38	11	...	0	0	
2	2013-08-21	26	426	...	0	0	
3	2014-10-02	26	11	...	0	0	
4	2014-01-19	94	173	...	0	0	

	AcceptedCmp5	AcceptedCmp1	AcceptedCmp2	Complain	Z_CostContact	\
0	0	0	0	0	0	3
1	0	0	0	0	0	3
2	0	0	0	0	0	3
3	0	0	0	0	0	3
4	0	0	0	0	0	3

	Z_Revenue	Response	Customer_For
0	11	1	971
1	11	0	125
2	11	0	472
3	11	0	65
4	11	0	321

[5 rows x 30 columns]

Now we will be exploring the unique values in the categorical features to get a clear idea of the data.

```
print("Total categories in the feature Marital_Status:\n",
df["Marital_Status"].value_counts(), "\n")
print("Total categories in the feature Education:\n",
df["Education"].value_counts())
```

Total categories in the feature Marital_Status:

Married	857
Together	573
Single	471
Divorced	232
Widow	76
Alone	3
Absurd	2
YOLO	2

Name: Marital_Status, dtype: int64

Total categories in the feature Education:

Graduation	1116
PhD	481
Master	365
2n Cycle	200
Basic	54

Name: Education, dtype: int64

In the next bit, I will be performing the following steps to engineer some new features:

- Extract the **"Age"** of a customer by the **"Year_Birth"** indicating the birth year of the respective person.
- Create another feature **"Spent"** indicating the total amount spent by the customer in various categories over the span of two years.
- Create another feature **"Living_With"** out of **"Marital_Status"** to extract the living situation of couples.
- Create a feature **"Children"** to indicate total children in a household that is, kids and teenagers.
- To get further clarity of household, Creating feature indicating **"Family_Size"**
- Create a feature **"Is_Parent"** to indicate parenthood status
- Lastly, I will create three categories in the **"Education"** by simplifying its value counts.
- Dropping some of the redundant features

```
print(df.columns)
```

```
Index(['ID', 'Year_Birth', 'Education', 'Marital_Status', 'Income',
      'Kidhome',
      'Teenhome', 'Dt_Customer', 'Recency', 'MntWines', 'MntFruits',
      'MntMeatProducts', 'MntFishProducts', 'MntSweetProducts',
      'MntGoldProds', 'NumDealsPurchases', 'NumWebPurchases',
```

```

        'NumCatalogPurchases', 'NumStorePurchases',
'NumWebVisitsMonth',
        'AcceptedCmp3', 'AcceptedCmp4', 'AcceptedCmp5', 'AcceptedCmp1',
        'AcceptedCmp2', 'Complain', 'Z_CostContact', 'Z_Revenue',
'Response',
        'Customer_For', 'Age'],
        dtype='object')

```

Feature Engineering

Age of customer today

```
df["Age"] = 2021 - df["Year_Birth"]
```

Total spendings on various items

```
df["Spent"] = df["MntWines"] + df["MntFruits"] + df["MntMeatProducts"]
+ df["MntFishProducts"] + df["MntSweetProducts"] + df["MntGoldProds"]
```

Deriving living situation by marital status "Alone"

```
df["Living_With"] = df["Marital_Status"].replace({"Married":
"Partner", "Together": "Partner", "Absurd": "Alone", "Widow": "Alone",
"YOLO": "Alone", "Divorced": "Alone", "Single": "Alone"})
```

Feature indicating total children living in the household

```
df["Children"] = df["Kidhome"] + df["Teenhome"]
```

Feature for total members in the household

```
df["Family_Size"] = df["Living_With"].replace({"Alone": 1, "Partner":
2}) + df["Children"]
```

Feature pertaining parenthood

```
df["Is_Parent"] = (df["Children"] > 0).astype(int)
```

Segmenting education levels into three groups

```
df["Education"] = df["Education"].replace({"Basic": "Undergraduate",
"2n Cycle": "Undergraduate", "Graduation": "Graduate", "Master":
"Postgraduate", "PhD": "Postgraduate"})
```

Renaming columns for clarity

```
df = df.rename(columns={"MntWines": "Wines", "MntFruits": "Fruits",
"MntMeatProducts": "Meat", "MntFishProducts": "Fish",
"MntSweetProducts": "Sweets", "MntGoldProds": "Gold"})
```

Dropping some redundant features

```
to_drop = ["Marital_Status", "Dt_Customer", "Z_CostContact",
"Z_Revenue", "Year_Birth", "ID"]
df = df.drop(to_drop, axis=1)
```

Display the resulting DataFrame

```
print(df.head())
```

```

      Education  Income  Kidhome  Teenhome  Recency  Wines  Fruits
Meat \

```

0	Graduate	58138.0	0	0	58	635	88
1	Graduate	46344.0	1	1	38	11	1
2	Graduate	71613.0	0	0	26	426	49
3	Graduate	26646.0	1	0	26	11	4
4	Postgraduate	58293.0	1	0	94	173	43

	Fish	Sweets	...	AcceptedCmp2	Complain	Response	Customer_For
Age \							
0	172	88	...	0	0	1	971
1	2	1	...	0	0	0	125
2	111	21	...	0	0	0	472
3	10	3	...	0	0	0	65
4	46	27	...	0	0	0	321

	Spent	Living_With	Children	Family_Size	Is_Parent
0	1617	Alone	0	1	0
1	27	Alone	2	3	1
2	776	Partner	0	2	0
3	53	Partner	1	3	1
4	422	Partner	1	3	1

[5 rows x 30 columns]

Now that we have some new features let's have a look at the data's stats.

```
df.describe()
```

	Income	Kidhome	Teenhome	Recency
Wines \				
count	2216.000000	2216.000000	2216.000000	2216.000000
mean	52247.251354	0.441787	0.505415	49.012635
std	25173.076661	0.536896	0.544181	28.948352
min	1730.000000	0.000000	0.000000	0.000000
25%	35303.000000	0.000000	0.000000	24.000000


```

50%      51381.500000      0.000000      0.000000      49.000000
174.500000
75%      68522.000000      1.000000      1.000000      74.000000
505.000000
max      66666.000000      2.000000      2.000000      99.000000
1493.000000

```

	Fruits	Meat	Fish	Sweets	Gold
... \					
count	2216.000000	2216.000000	2216.000000	2216.000000	2216.000000
...					
mean	26.356047	166.995939	37.637635	27.028881	43.965253
...					
std	39.793917	224.283273	54.752082	41.072046	51.815414
...					
min	0.000000	0.000000	0.000000	0.000000	0.000000
...					
25%	2.000000	16.000000	3.000000	1.000000	9.000000
...					
50%	8.000000	68.000000	12.000000	8.000000	24.500000
...					
75%	33.000000	232.250000	50.000000	33.000000	56.000000
...					
max	199.000000	1725.000000	259.000000	262.000000	321.000000
...					

	AcceptedCmp1	AcceptedCmp2	Complain	Response
Customer_For \				
count	2216.000000	2216.000000	2216.000000	2216.000000
2216.000000				
mean	0.064079	0.013538	0.009477	0.150271
512.006318				
std	0.244950	0.115588	0.096907	0.357417
232.469034				
min	0.000000	0.000000	0.000000	0.000000
0.000000				
25%	0.000000	0.000000	0.000000	0.000000
340.000000				
50%	0.000000	0.000000	0.000000	0.000000
513.000000				
75%	0.000000	0.000000	0.000000	0.000000
686.000000				
max	1.000000	1.000000	1.000000	1.000000
1063.000000				

	Age	Spent	Children	Family_Size	Is_Parent
count	2216.000000	2216.000000	2216.000000	2216.000000	2216.000000
mean	52.179603	607.075361	0.947202	2.592509	0.714350

std	11.985554	602.900476	0.749062	0.905722	0.451825
min	25.000000	5.000000	0.000000	1.000000	0.000000
25%	44.000000	69.000000	0.000000	2.000000	0.000000
50%	51.000000	396.500000	1.000000	3.000000	1.000000
75%	62.000000	1048.000000	1.000000	3.000000	1.000000
max	128.000000	2525.000000	3.000000	5.000000	1.000000

[8 rows x 28 columns]

The above stats show some discrepancies in mean Income and Age and max Income and age.

Do note that max-age is 128 years, As I calculated the age that would be today (i.e. 2021) and the data is old.

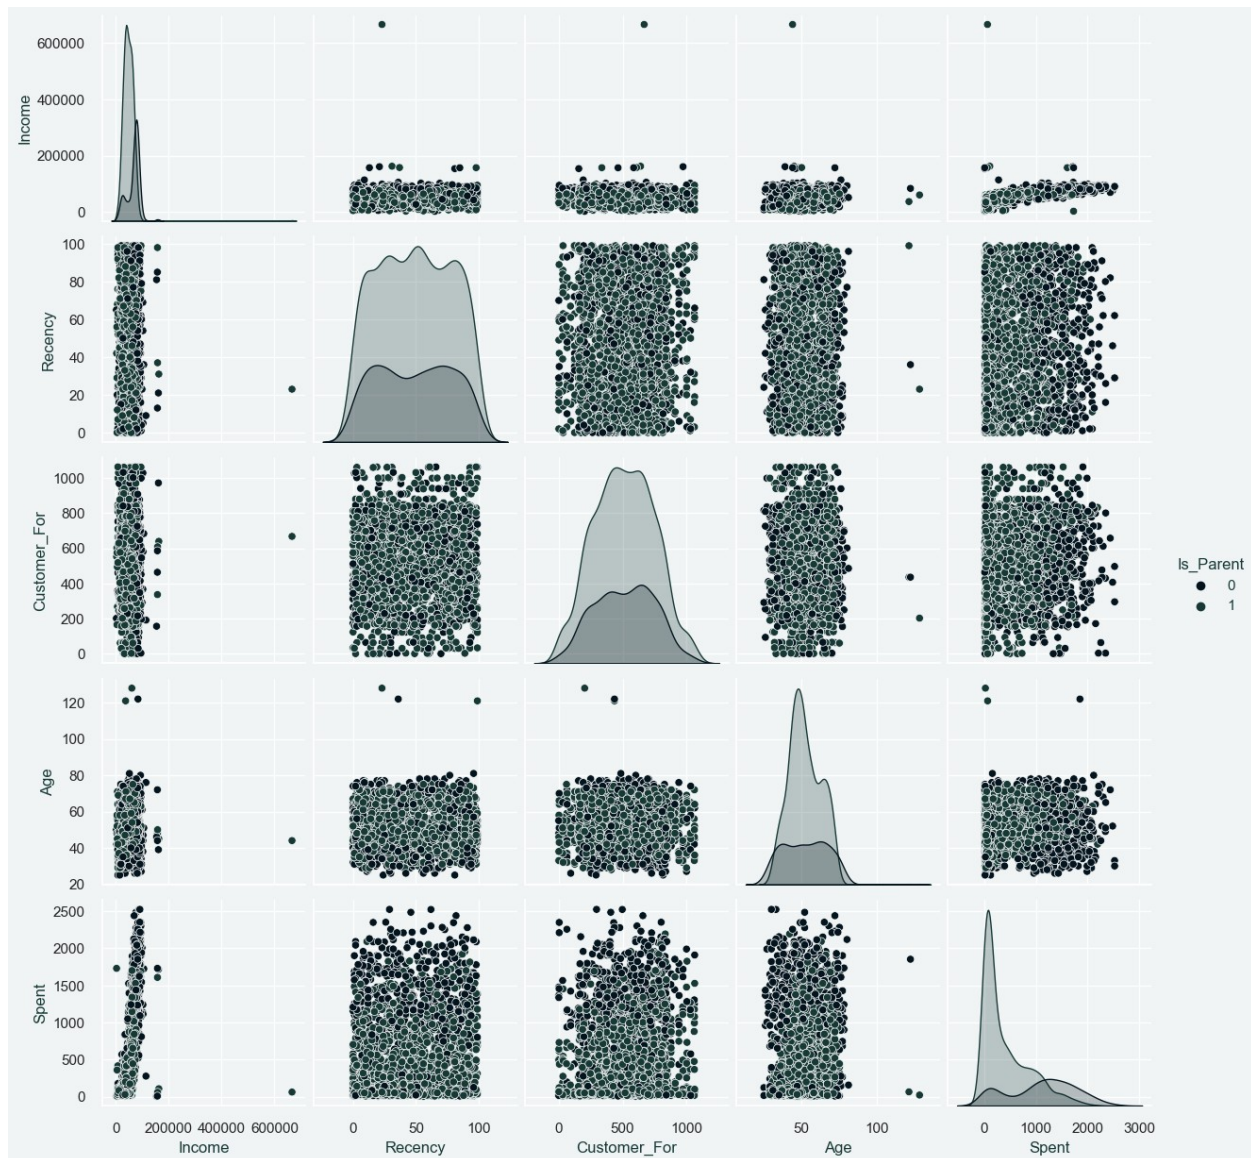
I must take a look at the broader view of the data. I will plot some of the selected features.

```
# Setting up colors preferences
sns.set(rc={"axes.facecolor": "#F0F4F4", "figure.facecolor":
"#F0F4F4", "text.color": "#183A37"})
pallet = ["#04151F", "#183A37", "#EFD6AC", "#C44900", "#432534",
"#CDB6C1"]
cmap = colors.ListedColormap(["#04151F", "#183A37", "#EFD6AC",
"#C44900", "#432534", "#CDB6C1"])

# Plotting selected features
to_plot = ["Income", "Recency", "Customer_For", "Age", "Spent",
"Is_Parent"]
print("Relative Plot Of Some Selected Features: A Data Subset")
plt.figure()
sns.pairplot(df[to_plot], hue="Is_Parent", palette=pallet)
plt.show()
```

Relative Plot Of Some Selected Features: A Data Subset

<Figure size 800x550 with 0 Axes>



Clearly, there are a few outliers in the Income and Age features. I will be deleting the outliers in the data.

```
# Dropping the outliers by setting a cap on Age and income
df = df[(df["Age"] < 90) & (df["Income"] < 600000)]
print("The total number of data-points after removing the outliers
are:", len(df))
```


The total number of data-points after removing the outliers are: 2212

Next, let us look at the correlation amongst the features. (Excluding the categorical attributes at this point)

```
# Correlation matrix
corrmat = df.corr()
```

```
plt.figure(figsize=(20, 20))
sns.heatmap(corrmat, annot=True, cmap=cmap, center=0)
```

<Axes: >



The data is quite clean and the new features have been included. I will proceed to the next step. That is, preprocessing the data.

DATA PREPROCESSING

In this section, I will be preprocessing the data to perform clustering operations.

The following steps are applied to preprocess the data:

- Label encoding the categorical features
- Scaling the features using the standard scaler
- Creating a subset dataframe for dimensionality reduction

```
# Get list of categorical variables
s = (df.dtypes == 'object')
object_cols = list(s[s].index)

print("Categorical variables in the dataset:", object_cols)

Categorical variables in the dataset: ['Education', 'Living_With']

# Label Encoding the object dtypes.
LE = LabelEncoder()
for i in object_cols:
    df[i] = df[[i]].apply(LE.fit_transform)

print("All features are now numerical")

All features are now numerical

# Creating a copy of df
ds = df.copy()
# Creating a subset of dataframe by dropping the features on deals
accepted and promotions
cols_del = ['AcceptedCmp3', 'AcceptedCmp4', 'AcceptedCmp5',
'AcceptedCmp1', 'AcceptedCmp2', 'Complain', 'Response']
ds = ds.drop(cols_del, axis=1)
# Scaling
scaler = StandardScaler()
scaler.fit(ds)
scaled_ds = pd.DataFrame(scaler.transform(ds), columns=ds.columns)
print("All features are now scaled")

All features are now scaled

# Scaled data to be used for reducing the dimensionality
print("Dataframe to be used for further modelling:")
scaled_ds.head()
```

Dataframe to be used for further modelling:

	Education	Income	Kidhome	Teenhome	Recency	Wines
Fruits \						
0	-0.893586	0.287105	-0.822754	-0.929699	0.310353	0.977660

```

1.552041
1 -0.893586 -0.260882 1.040021 0.908097 -0.380813 -0.872618 -
0.637461
2 -0.893586 0.913196 -0.822754 -0.929699 -0.795514 0.357935
0.570540
3 -0.893586 -1.176114 1.040021 -0.929699 -0.795514 -0.872618 -
0.561961
4 0.571657 0.294307 1.040021 -0.929699 1.554453 -0.392257
0.419540

```

```

      Meat      Fish      Sweets  ...  NumCatalogPurchases
NumStorePurchases \
0 1.690293 2.453472 1.483713 ...          2.503607 -
0.555814
1 -0.718230 -0.651004 -0.634019 ...          -0.571340 -
1.171160
2 -0.178542 1.339513 -0.147184 ...          -0.229679
1.290224
3 -0.655787 -0.504911 -0.585335 ...          -0.913000 -
0.555814
4 -0.218684 0.152508 -0.001133 ...          0.111982
0.059532

```

```

      NumWebVisitsMonth  Customer_For      Age      Spent  Living_With
Children \
0          0.692181          1.973583 1.018352 1.676245 -1.349603 -
1.264598
1         -0.132545         -1.665144 1.274785 -0.963297 -1.349603
1.404572
2         -0.544908         -0.172664 0.334530 0.280110 0.740959 -
1.264598
3          0.279818         -1.923210 -1.289547 -0.920135 0.740959
0.069987
4         -0.132545         -0.822130 -1.033114 -0.307562 0.740959
0.069987

```

```

      Family_Size  Is_Parent
0      -1.758359 -1.581139
1       0.449070 0.632456
2      -0.654644 -1.581139
3       0.449070 0.632456
4       0.449070 0.632456

```

```
[5 rows x 23 columns]
```

DIMENSIONALITY REDUCTION

In this problem, there are many factors on the basis of which the final classification will be done. These factors are basically attributes or features. The higher the number of features, the harder it is to work with it. Many of these features are correlated, and hence redundant. This is why I will be performing dimensionality reduction on the selected features before putting them through a classifier.

Dimensionality reduction is the process of reducing the number of random variables under consideration, by obtaining a set of principal variables.

Principal component analysis (PCA) is a technique for reducing the dimensionality of such datasets, increasing interpretability but at the same time minimizing information loss.

Steps in this section:

- Dimensionality reduction with PCA
- Plotting the reduced dataframe

Dimensionality reduction with PCA

For this project, I will be reducing the dimensions to 3.


```
#Initiating PCA to reduce dimention's aka features to 3
pca = PCA(n_components=3)
pca.fit(scaled_ds)
PCA_ds = pd.DataFrame(pca.transform(scaled_ds),
columns=(["col1", "col2", "col3"]))
PCA_ds.describe().T
```

	count	mean	std	min	25%	50%
75% \						
col1	2212.0	1.156399e-16	2.878377	-5.969394	-2.538494	-0.780421
	2.383290					
col2	2212.0	1.284887e-17	1.706839	-4.312196	-1.328316	-0.158123
	1.242289					
col3	2212.0	5.460771e-17	1.221956	-3.530416	-0.829067	-0.022692
	0.799895					
	max					
col1	7.444305					
col2	6.142721					
col3	6.611222					

```
# A 3D Projection Of Data In the Reduced Dimension
x = PCA_ds["col1"]
y = PCA_ds["col2"]
z = PCA_ds["col3"]

# To plot
fig = plt.figure(figsize=(10, 8))
```

```
ax = fig.add_subplot(111, projection="3d")
ax.scatter(x, y, z, c="#183A37", marker="o")
ax.set_title("A 3D Projection Of Data In The Reduced Dimension")
plt.show()
```



CLUSTERING

Now that I have reduced the attributes to three dimensions, I will be performing clustering via Agglomerative clustering. Agglomerative clustering is a hierarchical clustering method. It involves merging examples until the desired number of clusters is achieved.

Steps involved in the Clustering

- Elbow Method to determine the number of clusters to be formed
- Clustering via Agglomerative Clustering
- Examining the clusters formed via scatter plot

```
# Quick examination of elbow method to find numbers of clusters to make.
```

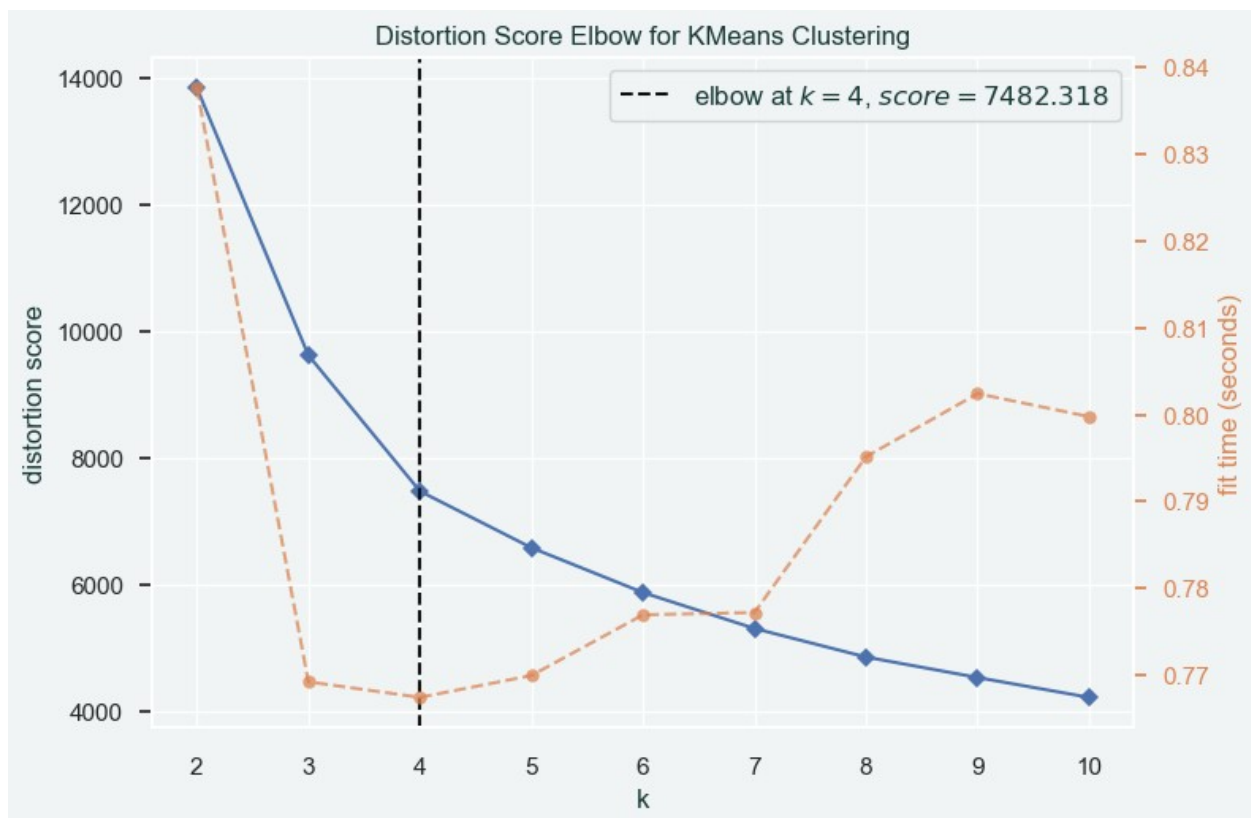
```
print('Elbow Method to determine the number of clusters to be formed:')
```

```
Elbow_M = KElbowVisualizer(KMeans(), k=10)
```

```
Elbow_M.fit(PCA_ds)
```

```
Elbow_M.show()
```

Elbow Method to determine the number of clusters to be formed:



```
<Axes: title={'center': 'Distortion Score Elbow for KMeans Clustering'}, xlabel='k', ylabel='distortion score'>
```


The above cell indicates that four will be an optimal number of clusters for this data. Next, we will be fitting the Agglomerative Clustering Model to get the final clusters.

```
#Initiating the Agglomerative Clustering model
AC = AgglomerativeClustering(n_clusters=4)
# fit model and predict clusters
yhat_AC = AC.fit_predict(PCA_ds)
PCA_ds["Clusters"] = yhat_AC
#Adding the Clusters feature to the original dataframe.
df["Clusters"]= yhat_AC
```

To examine the clusters formed let's have a look at the 3-D distribution of the clusters.

```
# Define the color palette using the provided color code
cluster_colors = ["#04151F", "#183A37", "#EFD6AC", "#C44900",
"#432534", "#CDB6C1"]
cmap = ListedColormap(cluster_colors)

# Plotting the clusters
fig = plt.figure(figsize=(10, 8))
ax = fig.add_subplot(111, projection='3d')
ax.scatter(x, y, z, s=40, c=PCA_ds["Clusters"], marker='o', cmap=cmap)
ax.set_title("The Plot Of The Clusters")
plt.show()
```



EVALUATING MODELS

Since this is an unsupervised clustering, We do not have a tagged feature to evaluate or score our model. The purpose of this section is to study the patterns in the clusters formed and determine the nature of the clusters' patterns.

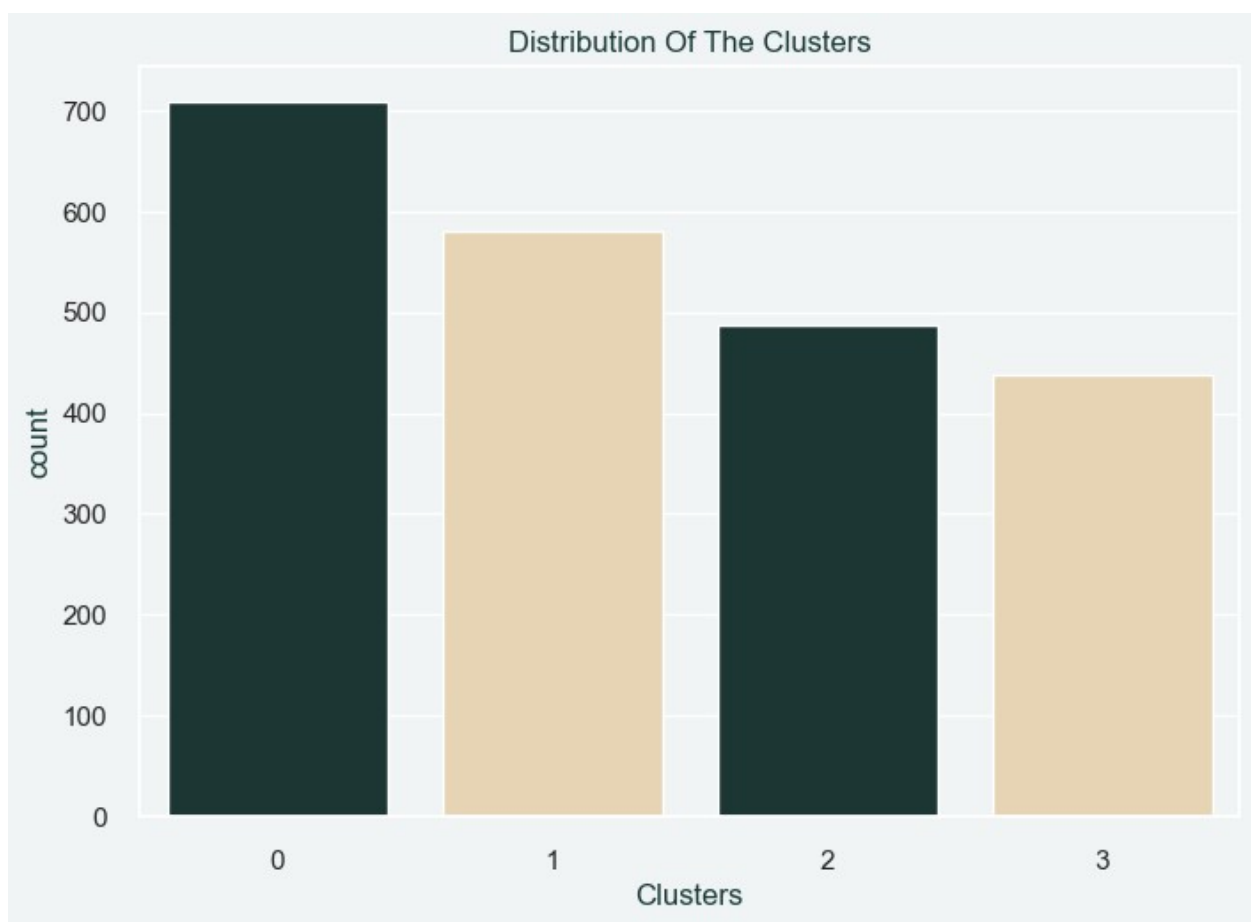
For that, we will be having a look at the data in light of clusters via exploratory data analysis and drawing conclusions.

Firstly, let us have a look at the group distribution of clustering

```
color1 = "#183A37"
color2 = "#EFD6AC"


# Generate a color palette by interpolating between the two colors
n_colors = len(pal)
custom_palette = sns.color_palette([color1, color2], n_colors)

# Plotting countplot of clusters
pl = sns.countplot(x=df["Clusters"], palette=custom_palette)
pl.set_title("Distribution Of The Clusters")
plt.show()
```



The clusters seem to be fairly distributed.

```
pl = sns.scatterplot(data=df, x=df["Spent"], y=df["Income"],
hue=df["Clusters"], palette=pal)
pl.set_title("Cluster's Profile Based On Income And Spending")
plt.legend()
plt.show()
```

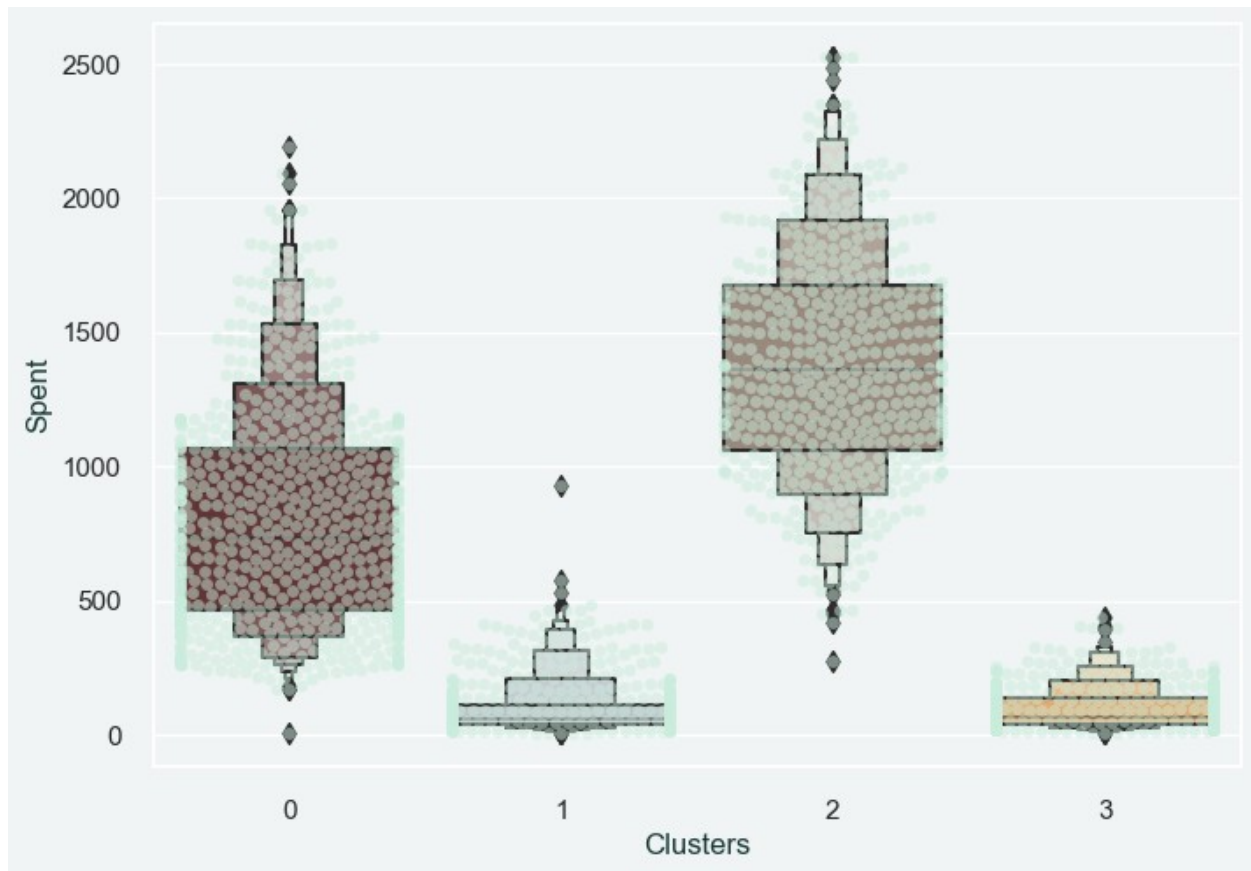


Income vs spending plot shows the clusters pattern

- group 0: high spending & average income
- group 1: high spending & high income
- group 2: low spending & low income
- group 3: high spending & low income

Next, I will be looking at the detailed distribution of clusters as per the various products in the data. Namely: Wines, Fruits, Meat, Fish, Sweets and Gold

```
plt.figure()
pl = sns.swarmplot(x=df["Clusters"], y=df["Spent"], color="#CBEDDD",
alpha=0.5)
pl = sns.boxenplot(x=df["Clusters"], y=df["Spent"], palette=pal)
plt.show()
```




From the above plot, it can be clearly seen that cluster 1 is our biggest set of customers closely followed by cluster 0. We can explore what each cluster is spending on for the targeted marketing strategies.

Let us next explore how did our campaigns do in the past.

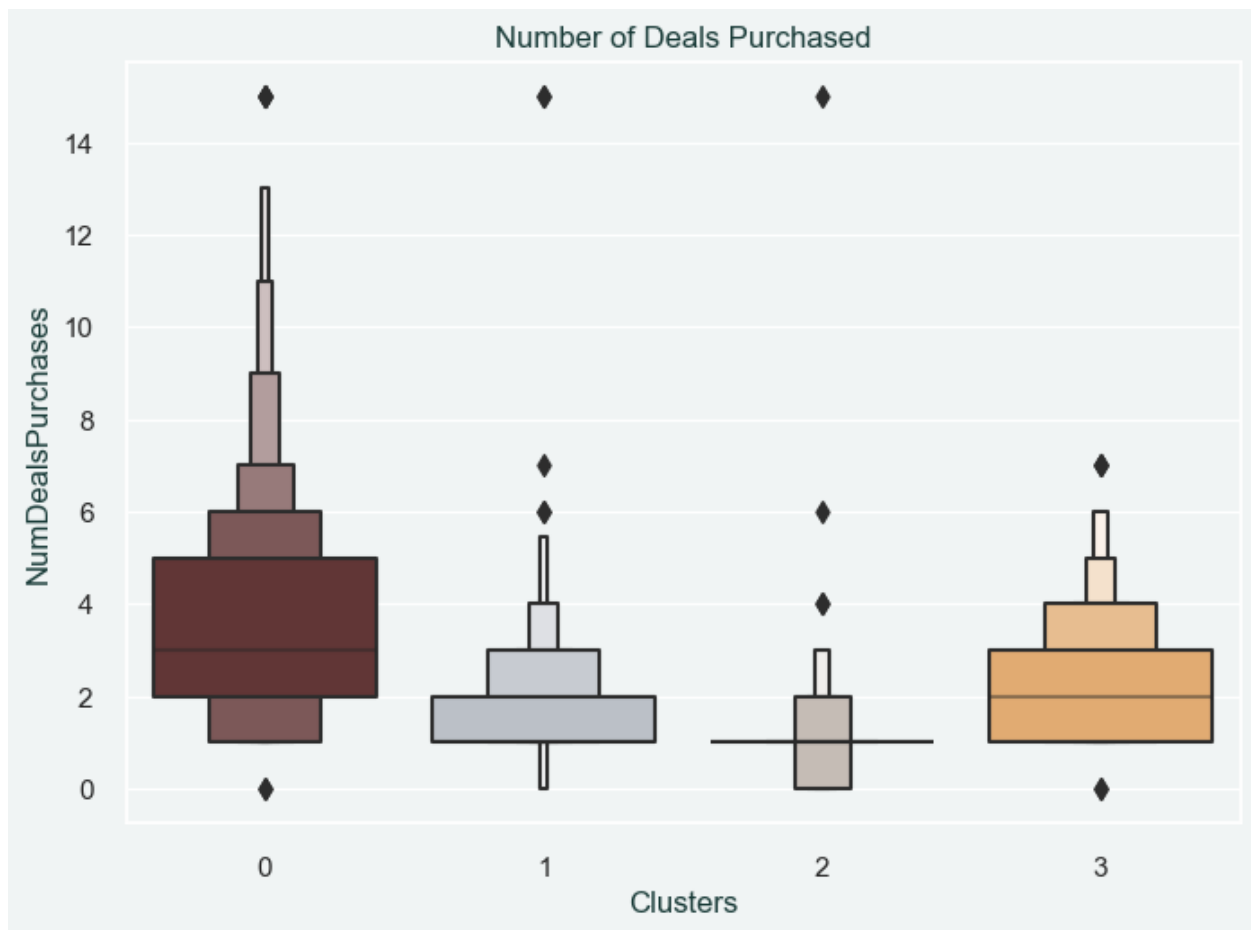
```
# Define the custom color palette with four colors
custom_palette = ["#183A37", "#EFD6AC", "#C44900", "#04151F"]

# Plotting count of total campaign accepted with the custom color
palette
plt.figure()
pl = sns.countplot(x=df["Total_Promos"], hue=df["Clusters"],
palette=custom_palette)
pl.set_title("Count Of Promotion Accepted")
pl.set_xlabel("Number Of Total Accepted Promotions")
plt.show()
```



There has not been an overwhelming response to the campaigns so far. Very few participants overall. Moreover, no one part take in all 5 of them. Perhaps better-targeted and well-planned campaigns are required to boost sales.

```
# Plotting the number of deals purchased
plt.figure()
pl = sns.boxenplot(y=df["NumDealsPurchases"], x=df["Clusters"],
palette=pal)
pl.set_title("Number of Deals Purchased")
plt.show()
```




Unlike campaigns, the deals offered did well. It has best outcome with cluster 0 and cluster 3. However, our star customers cluster 1 are not much into the deals. Nothing seems to attract cluster 2 overwhelmingly


```
# For more details on the purchasing style
Places = ["NumWebPurchases", "NumCatalogPurchases",
          "NumStorePurchases", "NumWebVisitsMonth"]

for i in Places:
    plt.figure(figsize=(6, 6))
    sns.jointplot(x=df[i], y=df["Spent"], hue=df["Clusters"],
                  palette=pal)
    plt.show()


<Figure size 600x600 with 0 Axes>
```


<Figure size 600x600 with 0 Axes>



<Figure size 600x600 with 0 Axes>



<Figure size 600x600 with 0 Axes>



PROFILING

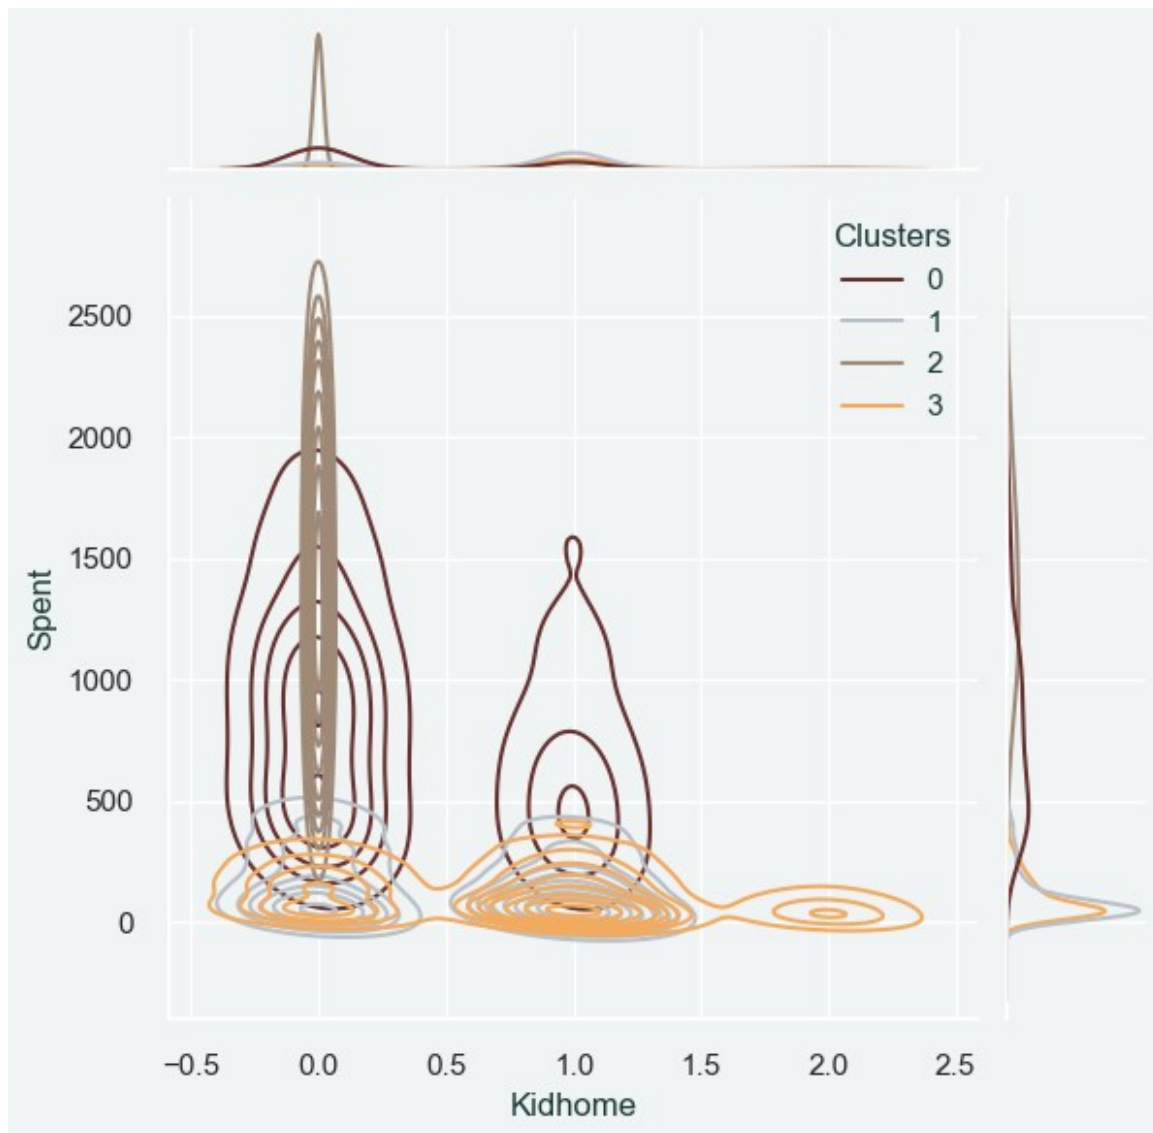
Now that we have formed the clusters and looked at their purchasing habits. Let us see who all are there in these clusters. For that, we will be profiling the clusters formed and come to a conclusion about who is our star customer and who needs more attention from the retail store's marketing team.

To decide that I will be plotting some of the features that are indicative of the customer's personal traits in light of the cluster they are in. On the basis of the outcomes, I will be arriving at the conclusions.


```
Personal = ["Kidhome", "Teenhome", "Customer_For", "Age", "Children",
"Family_Size", "Is_Parent", "Education", "Living_With"]
```

```
for i in Personal:
    plt.figure(figsize=(8, 4)) # Adjust the height (second value) as
    # needed
    sns.jointplot(x=df[i], y=df["Spent"], hue=df["Clusters"],
    kind="kde", palette=pal)
    plt.show()
```

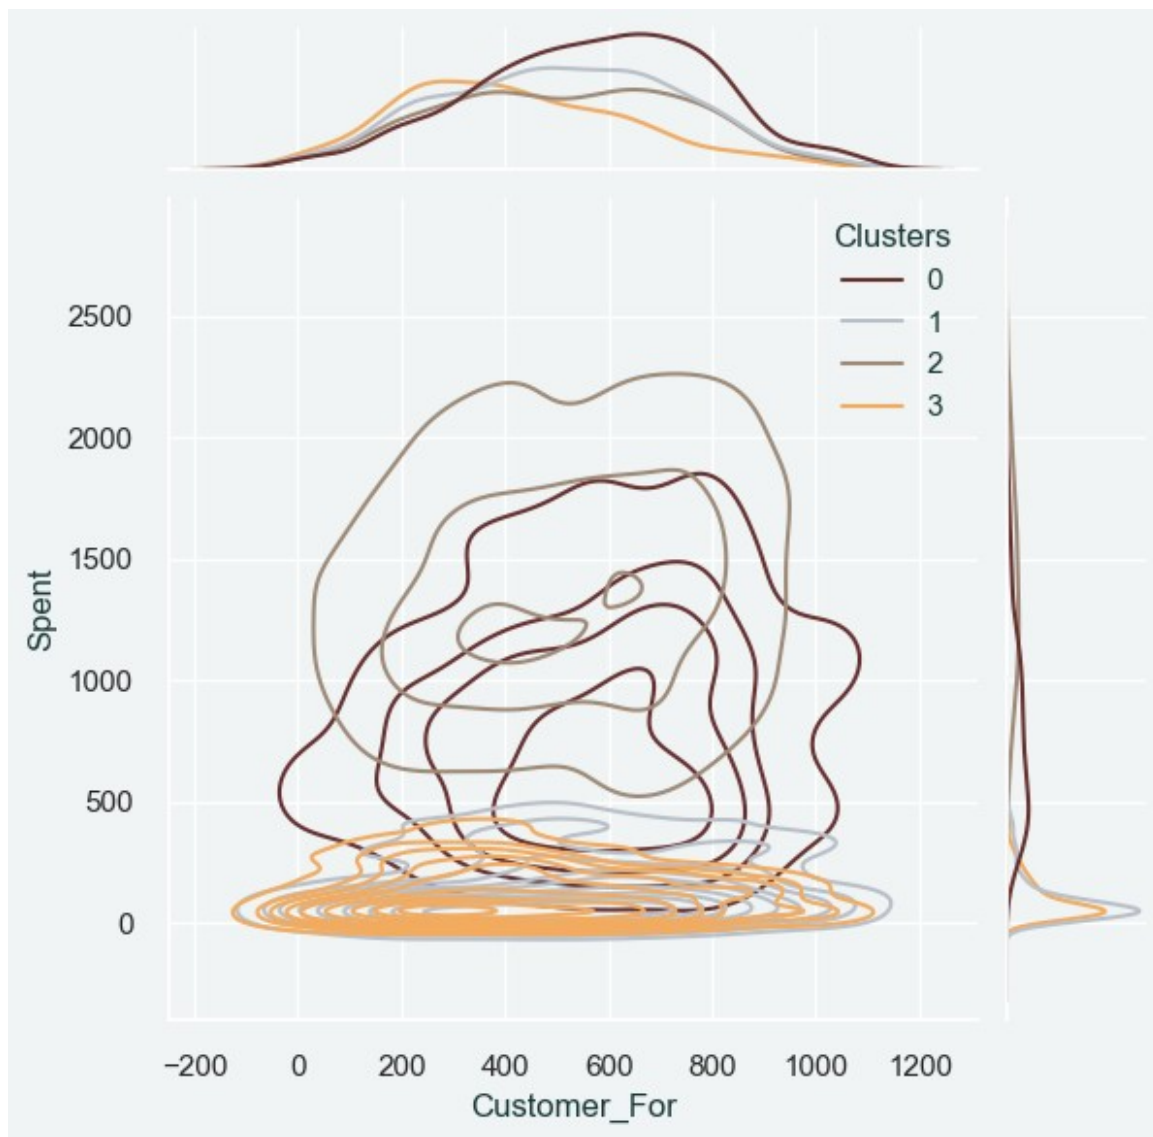
<Figure size 800x400 with 0 Axes>



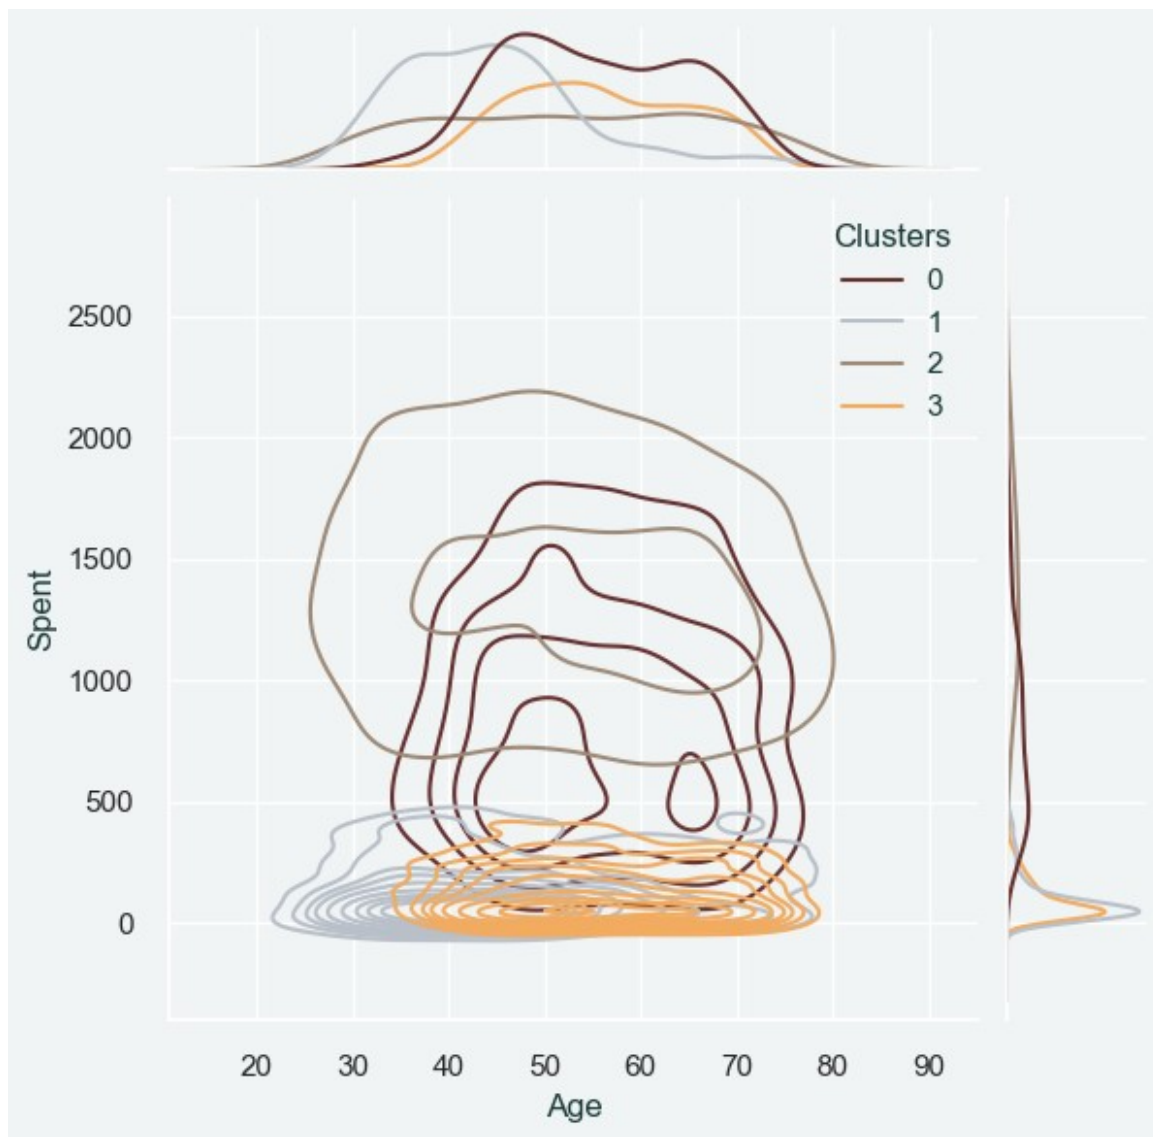
<Figure size 800x400 with 0 Axes>




<Figure size 800x400 with 0 Axes>




<Figure size 800x400 with 0 Axes>



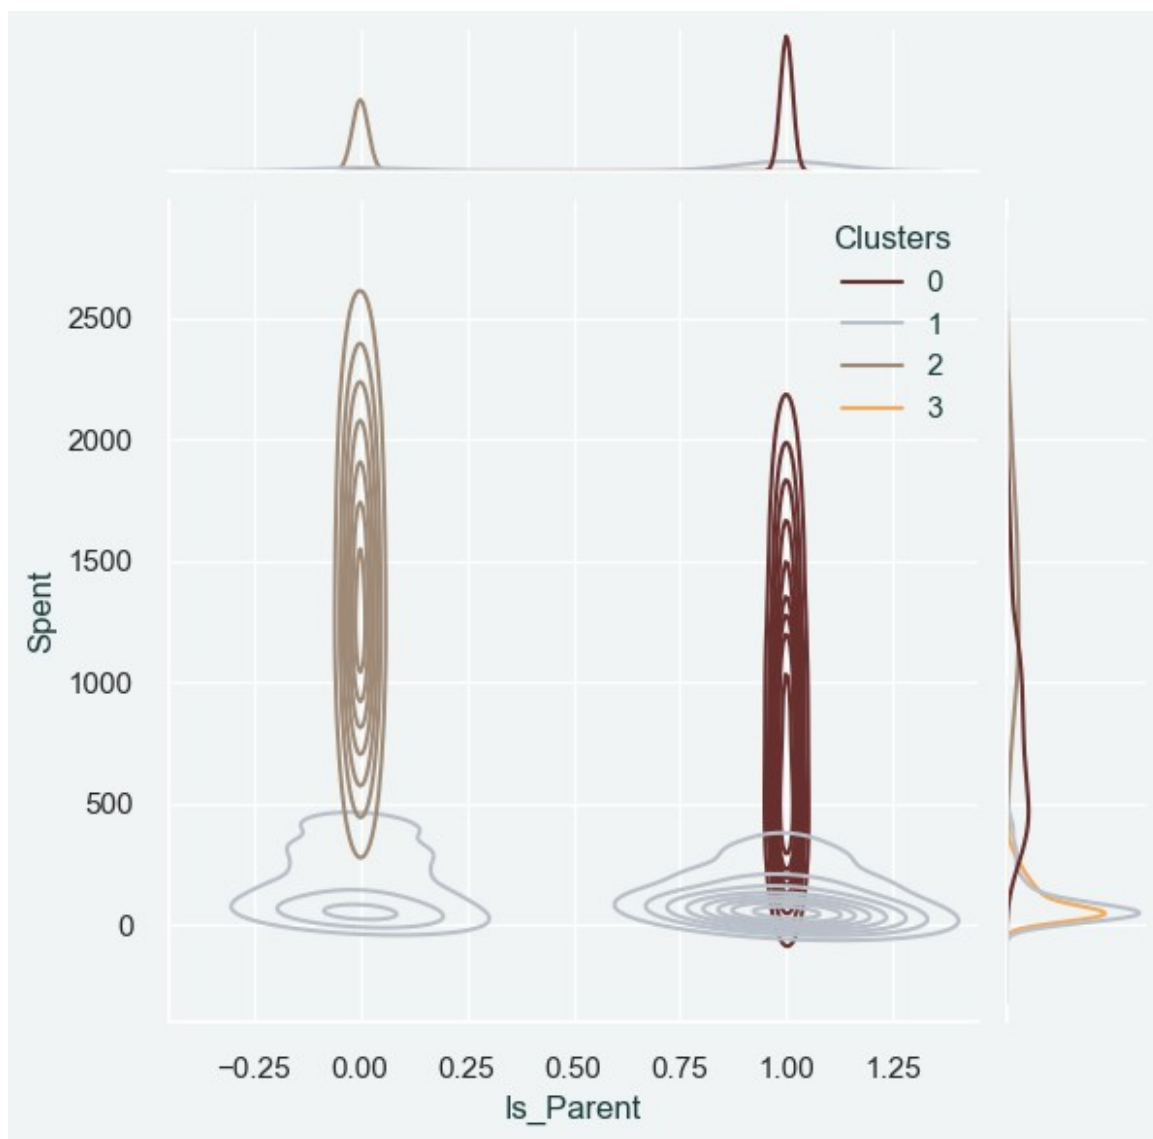
<Figure size 800x400 with 0 Axes>




<Figure size 800x400 with 0 Axes>




<Figure size 800x400 with 0 Axes>



<Figure size 800x400 with 0 Axes>



<Figure size 800x400 with 0 Axes>



Points to be noted:

The following information can be deduced about the customers in different clusters.

CONCLUSION

In this project, I performed unsupervised clustering. I did use dimensionality reduction followed by agglomerative clustering. I came up with 4 clusters and further used them in profiling customers in clusters according to their family structures and income/spending. This can be used in planning better marketing strategies.

If you liked this Notebook, please do upvote.

If you have any questions, feel free to comment!

END

